

This is a general purpose image classifier that can be applied to datasets that have the data organized as shown below. To use this code select Run all. The code will prompt the user to provide the required input data. It will take only a minute or so to enter the data and your model will be up and training. I have used this code on numerous data sets with excellent results. The model will initially run for 10 epochs. After 10 epochs you are prompted to enter H to halt training or you can enter an integer that specifies how many more epochs to run. After those epochs complete you will be again prompted to continue or halt training. The code shows the % decreasee of the validation loss at the end of each epoch. This is a good guide to use to decide to continue or halt training. As you near the bottom of the loss function the % decrease in loss becomes smaller. I typically halt training when the % decrease falls below 5%. However sometimes it is best to continue training until the loss increases. The code then automatically lowers the learning rate. This often enables the model to again start decreasing the loss. Your saved model is always returned with the weights set to those for the epoch with the lowest loss.

data directory

- train directory
 - class 0
 - image 0
 - image 1
 -
 - image m
 - class 1
 - image 0
 - image 1
 -
 - image n
 -
 - class m
 - image 0
 - image 1
 -
 - image r
- valid directory (optional if there is no validation directory validation data is created by partition of train data)
 - class 0
 - image 0
 - image 1
 -
 - image x
 - class 1
 - image 0

- image 1
-
- image y
-
- class m
 - image 0
 - image 1
 -
 - image z
- test directory (optional if there is no test directory test data is created by partition of train data)
 - class 0
 - image 0
 - image 1
 -
 - image a
 - class 1
 - image 0
 - image 1
 -
 - image b
 -
 - class m
 - image 0
 - image 1
 -
 - image c

1. Import Needed Modules

2 Define a function to print text in color

3. Define a function to plot the number of images in dataset classes

4 Define a function to return information on the dataset

5. Read in images and create a dataframe of image paths and class labels

6. Define a function that trims a dataset for the max number of class images

7. Balance the training set using augmentation
8. Create train, test and validation generators
9. Create a function to show Training Image Samples
- 10 Create a function to calculate the F1 score metric
11. Create the Model
12. Create a custom Keras callback to continue or halt training
13. Define a function to plot the training data
14. Define a function save the training data to a csv file
15. Define a function to make predictions on the test set
16. Make predictions on test set, create Confusion Matrix and Classification Report
17. Define a function to save the trained model
- 18 Define the run function which runs the classifier
- 19 Define the code to enable to initiate the run function
20. Evaluate Model Performance

Recently Kaggle updated its docker file to load tensorflow 2.11.0. where before it loaded version 2.9.2 . However version 2.10.0 and above have a bug with respect to saving EfficientNet models. In my code if you select a medium or large model, the result is an EfficientNetB0 or B3 model is generated. So when you try to save the model you will receive a message saying the model can not be saved. If you want to save your model, uncomment the code line below which will install tensorflow version 2.9.2.

```
In [1]: #! pip install tensorflow==2.9.2
```

Import required modules

```
In [2]: import pandas as pd
import numpy as np
from numpy import random
import os
os.environ['TF_CPP_MIN_LOG_LEVEL'] = '2'
import time
import matplotlib
import matplotlib.pyplot as plt
import cv2
import seaborn as sns
sns.set_style('darkgrid')
import shutil
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.model_selection import train_test_split
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.layers import Dense, Activation, Dropout, Conv2D, MaxPooling
from tensorflow.keras.optimizers import Adam, Adamax
from tensorflow.keras.metrics import categorical_crossentropy
from tensorflow.keras import regularizers
from tensorflow.keras.models import Model
from tensorflow.keras import backend as K
import albumentations as A
import time
from tqdm import tqdm
from sklearn.metrics import f1_score
from IPython.display import YouTubeVideo
import sys
if not sys.warnoptions:
    import warnings
    warnings.simplefilter("ignore")
pd.set_option('display.max_columns', None) # or 1000
pd.set_option('display.max_rows', None) # or 1000
pd.set_option('display.max_colwidth', None) # or 199
```

```
/opt/conda/lib/python3.10/site-packages/scipy/__init__.py:146: UserWarning: A NumPy version >=1.16.5 and <1.23.0 is required for this version of SciPy (detected version 1.24.3)
warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}"
```

```
In [3]: # Seed Everything to reproduce results for future use cases
def seed_everything(seed=42):
    # Seed value for TensorFlow
    tf.random.set_seed(seed)

    # Seed value for NumPy
    np.random.seed(seed)

    # Seed value for Python's random library
    random.seed(seed)

    # Force TensorFlow to use single thread
    # Multiple threads are a potential source of non-reproducible results.
    session_conf = tf.compat.v1.ConfigProto(
        intra_op_parallelism_threads=1,
        inter_op_parallelism_threads=1
    )

    # Make sure that TensorFlow uses a deterministic operation wherever possible
    tf.compat.v1.set_random_seed(seed)

    sess = tf.compat.v1.Session(graph=tf.compat.v1.get_default_graph(), config=
        tf.compat.v1.keras.backend.set_session(sess))

seed_everything()
```

Define a function to print text in specified rgb foreground and background colors

Add some PZAZZ to your printed output with this function

form of the call is: `print_in_color(txt_msg, fore_tuple, back_tuple)` where:

- `txt_msg` is the string to be printed out
- `fore_tuple` is tuple of the form `(r,g,b)` specifying the foreground color of the text
- `back_tuple` is tuple of the form `(r,g,b)` specifying the background color of the text

```
In [4]: # for kaggle use (0,0,0) and (255,255,255)
def print_in_color(txt_msg,fore_tuple= (0,255,255) ,back_tuple=(0,0,0)):
    #prints the text_msg in the foreground color specified by fore_tuple with t
    #text_msg is the text, fore_tuple is foregroud color tuple (r,g,b), back_t
    rf,gf,bf=fore_tuple
    rb,gb,bb=back_tuple
    msg='{0}' + txt_msg
    mat='\33[38;2;' + str(rf) +';' + str(gf) + ';' + str(bf) + ';48;2;' + str(rb
```

```
print(msg .format(mat))
print('\33[0m', end='') # returns default print color to back to black
```

Define a function that plots value counts for a column in a dataframe

```
In [5]: def plot_label_count (df, plot_title):
    column='labels'
    xaxis_label='CLASS'
    yaxis_label='IMAGE COUNT'
    vcounts=df[column].value_counts()
    labels=vcounts.keys().tolist()
    values=vcounts.tolist()
    lcount=len(labels)
    if lcount>60:
        print_in_color('The number of labels is >55, no plot will be produced')
    else:
        width=lcount * 4
        width=np.min([width, 20])
        plt.figure(figsize=(width,5))
        form = {'family': 'serif', 'color': 'blue', 'size': 25}
        plt.bar(labels, values)
        plt.title(plot_title, fontsize= 24, color='blue')
        plt.xticks(rotation=90, fontsize=18)
        plt.yticks(fontsize=18)
        plt.xlabel(xaxis_label, fontdict=form)
        plt.ylabel(yaxis_label, fontdict=form)
        if lcount >=8:
            rotation='vertical'
        else:
            rotation='horizontal'
        for i in range(lcount):
            plt.text(i, values[i]/2, str(values[i]), fontsize=12, rotation=rotation)
        plt.show()
```

Define a function to return information of the training data

- train_dir is a string specifying the full path to the directory holding the training images

The function determines the total number of training images, the class with the most image files

and their number and the class with the least image files and their number. This information is

provided to the user to enable selection of various program parameters

```
In [6]: def check_dataset_size(train_dir):
    classes=sorted(os.listdir(train_dir))
    ftotal=0
    flargest=0
```

```
fsmallest=1000000000
for klass in classes:
    classpath=os.path.join(train_dir,klass)
    if os.path.isdir(classpath):
        plist=os.listdir(classpath)
        fcount=len(plist)
        if fcount>flargest:
            flargest=fcount
            maxclass=klass
        if fcount < fsmallest:
            fsmallest=fcount
            minclass=klass
    ftotal += fcount
return ftotal, flargest, maxclass, fsmallest, minclass
```

Read in data and create train, test and validation data frames

```
In [7]: def make_dataframes(train_dir,test_dir, val_dir, limiter, min_images):
    bad_images=[]
    skip_classes=[]
    # check what directories exist
    if test_dir == None and val_dir==None:
        dirlist=[train_dir]
        names = ['train']
    elif test_dir == None:
        dirlist=[train_dir, val_dir]
        names=['train', 'valid']
    elif val_dir == None:
        dirlist=[train_dir, test_dir]
        names=['train', 'test']
    else:
        dirlist=[train_dir, test_dir, val_dir]
        names=['train','test', 'valid']
    ht=0 # set initial value of height counter
    wt=0 # set initial value of width counter
    total_good_files=0 # set initial value of total number of good image files
    zipdir=zip(names, dirlist)
    for name,d in zipdir: #iterate through the names and directories
        filepaths=[] # initialize list of filepaths
        labels=[] # initialize list of class labels
        classlist=sorted(os.listdir(d)) # get a list of all the classes in a directory
        for klass in classlist: # iterate through the list of classes
            msg=f'processing images in {name} directory for class {klass}'
            print(msg, '\r', end='')
            good_file_count=0 # initialize the good_file count for this class
            classpath=os.path.join(d, klass) # define the full path to the class
            if os.path.isdir(classpath): # ensure we are working with a directory
                plist=sorted(os.listdir(classpath)) # make a list of all the files
                if len(plist)< min_images:
                    skip_classes.append(klass)
            else:
                if klass not in skip_classes:
                    if limiter != None: # check if a limiter value was specified
                        if len(plist)<limiter: # if there are more files than the limiter
                            plist=np.random.choice(plist, limiter, replace=False)
```

```

        for f in flist:
            fpath=os.path.join(classpath,f) # create the full pa
            index=f.rfind('.')
            ext=f[index+1:].lower() # the the file's extension
            if ext not in ['jpg', 'jpeg', 'tiff', 'png', 'bmp']
                bad_images.append(fpath) # if not a proper exten
            else:
                try: # check if image files are defective if so
                    img=cv2.imread(fpath)
                    h=img.shape[0]
                    w=img.shape[1]
                    ht +=h # add images height and width to the
                    wt += w
                    good_file_count +=1
                    total_good_files +=1
                    filepaths.append(fpath) # append the filepat
                    labels.append(klass) # append the file's cla

                except:
                    bad_images.append(fpath) # if the image file
                    print('')

        Fseries=pd.Series(filepaths, name='filepaths') # make a pandas series fo
        Lseries=pd.Series(labels, name='labels')
        df=pd.concat([Fseries, Lseries], axis=1) # make a dataframe with columns
        # depending on which directory we are iterating through create dataframe
        if name == 'valid':
            valid_df=df
        elif name == 'test':
            test_df=df
        else:
            if test_dir == None and val_dir == None: # create train_df, test_df
                pdf=df
                train_df, dummy_df=train_test_split(pdf, train_size=.8, shuffle=True)
                valid_df, test_df=train_test_split(dummy_df, train_size=.5, shuf
            elif test_dir == None: # create just a train_df and test_df
                pdf=df
                train_df,test_df=train_test_split(pdf, train_size=.8, shuffle=True)
            elif val_dir == None: # create a train_df and a valid_df
                pdf=df
                train_df,valid_df=train_test_split(pdf, train_size=.8, shuffle=True)
            else:
                train_df= df # test and valid dataframes exists so train_df is j
        classes=sorted(train_df['labels'].unique())
        class_count=len(classes)
        # calculate the average image height and with
        have=int(ht/total_good_files)
        wave=int(wt/total_good_files)
        aspect_ratio=have/wave
        print('number of classes in processed dataset= ', class_count)
        if len(skip_classes)>0:
            msg=f'listed below are classes not included in the dataset because it ha
            print_in_color(msg)
            for klass in skip_classes:
                print (klass)

        counts=list(train_df['labels'].value_counts())
        print('the maximum files in any class in train_df is ', max(counts), ' the
        print('train_df length: ', len(train_df), ' test_df length: ', len(test_df))
        if len(bad_images) == 0:
            print_in_color('All image files were properly processed and used in the

```

```

else:
    print_in_color(f'the are {len(bad_images)} bad image files and {total_good_images} good ones')
    for f in bad_images:
        print(f)
plot_title='Images per Label in train set'
plot_label_count (train_df, plot_title)
return train_df, test_df, valid_df, classes, class_count, max(counts), min(counts)

```

Define a function that trims the classes in a dataframe

- df is the dataframe
- max_samples is an integer specifies the maximum number of images a class can have
- min_samples is an integer specifying the minimum number of images a class must have to be included in the trimmed dataframe
- column is a string specifying the column name in the dataframe holding the class labels

```
In [8]: def trim(df, max_samples, min_samples, column):
    # column specifies which column of the dataframe to use, typically this is the class label
    # df is typically train_df
    df=df.copy()
    classes=df[column].unique() # get the classes in df
    class_count=len(classes)
    length=len(df)
    print ('dataframe initially is of length ',length, ' with ', class_count, ' classes')
    groups=df.groupby(column) # creates a set of dataframes that only contain images for a specific class
    trimmed_df = pd.DataFrame(columns = df.columns) # create an empty dataframe
    for label in df[column].unique(): # iterate through each class label
        group=groups.get_group(label) # get the dataframe associate with the label
        count=len(group) # determine how many files are in the dataframe
        if count > max_samples: # if there more files in the dataframe sample it
            sampled_group=group.sample(n=max_samples, random_state=123, axis=0)
            trimmed_df=pd.concat([trimmed_df, sampled_group], axis=0) # add the sampled group to the trimmed dataframe
        else:
            if count>=min_samples: # if the dataframe has more than the minimum required samples
                sampled_group=group
            trimmed_df=pd.concat([trimmed_df, sampled_group], axis=0)
    print('after trimming, the maximum samples in any class is now ',max_samples)
    classes=trimmed_df[column].unique()# return this in case some classes have less than max samples
    class_count=len(classes) # return this in case some classes have less than max samples
    length=len(trimmed_df)
    print ('the trimmed dataframe now is of length ',length, ' with ', class_count, ' classes')
    return trimmed_df, classes, class_count
```

Expand train_df rows with augmented images so each class has n samples

- df is the dataframe
- n is an integer specifying the number of images desired for each class

- column is a string specifying the column name of the dataframe that contains the class labels
- working_dir is a string specifying the path where augmented images will be stored
- img_size is a tuple (height,width) specifying the image size of the augmented images created

```
In [9]: def balance(df, n, column, working_dir, img_size):
    def get_augmented_image(image): # given an image this function returns an augmented image
        width=int(image.shape[1]*.8)
        height=int(image.shape[0]*.8)
        transform= A.Compose([
            A.HorizontalFlip(p=.5),
            A.Rotate(limit=30, p=.25),
            A.RandomBrightnessContrast(p=.5),
            A.RandomGamma(p=.5),
            A.RandomCrop(width=width, height=height, p=.25) ])
        return transform(image=image)['image']
    def dummy(image):
        return image

    df=df.copy()
    print('Initial length of dataframe is ', len(df))
    aug_dir=os.path.join(working_dir, 'aug')# directory to store augmented images
    if os.path.isdir(aug_dir):# start with an empty directory
        shutil.rmtree(aug_dir)
    os.mkdir(aug_dir)
    for label in df[column].unique():
        dir_path=os.path.join(aug_dir,label)
        os.mkdir(dir_path) # make class directories within aug directory
    # create and store the augmented images
    total=0
    groups=df.groupby(column) # group by class
    for label in df[column].unique(): # for every class
        msg=f'augmenting images in train set for class {label}'
        print(msg, '\r', end='')
        group=groups.get_group(label) # a dataframe holding only rows with the target class
        sample_count=len(group) # determine how many samples there are in this group
        if sample_count<n: # if the class has less than target number of images
            aug_img_count=0
            delta=n - sample_count # number of augmented images to create
            target_dir=os.path.join(aug_dir, label) # define where to write the augmented images
            desc=f'augmenting class {label:25s}'
            for i in range(delta):
                j= i % sample_count # need this because we may have to go through the group multiple times
                img_path=group['filepaths'].iloc[j]
                img=cv2.imread(img_path)
                img=get_augmented_image(img)
                fname=os.path.basename(img_path)
                fname='aug' +str(i) + '-' +fname
                dest_path=os.path.join(target_dir, fname)
                cv2.imwrite(dest_path, img)
                aug_img_count +=1
            total +=aug_img_count
        print('')
        print('Total Augmented images created= ', total)
    # create aug_df and merge with train_df to create composite training set ndf
    aug_fpaths=[ ]
```

```
aug_labels=[]
classlist=sorted(os.listdir(aug_dir))
for klass in classlist:
    classpath=os.path.join(aug_dir, klass)
    flist=sorted(os.listdir(classpath))
    for f in flist:
        fpath=os.path.join(classpath,f)
        aug_fpaths.append(fpath)
        aug_labels.append(klass)
Fseries=pd.Series(aug_fpaths, name='filepaths')
Lseries=pd.Series(aug_labels, name='labels')
aug_df=pd.concat([Fseries, Lseries], axis=1)
df=pd.concat([df,aug_df], axis=0).reset_index(drop=True)
print('Length of augmented dataframe is now ', len(df))
return df
```

Create the train_gen, test_gen and valid_gen

- batch_size is an integer specifying the generator batch size
- ycol is a string specifying the dataframe column containing the class labels
- train_df is the dataframe containing the training images
- test_df is the dataframe containing the test images
- valid_df is the dataframe containing the validation images
- image_size is a tuple (height, width)

```
In [10]: def make_gens(batch_size, ycol, train_df, test_df, valid_df, img_size, modnum):
    if modnum == 6:
        gen=ImageDataGenerator(preprocessing_function= tf.keras.applications.inception_v3.preprocess_input)
    else:
        gen=ImageDataGenerator()
    msg='{0:70s} for train generator'.format(' ')
    print(msg, '\r', end='') # prints over on the same line
    train_gen=gen.flow_from_dataframe(train_df, x_col='filepaths', y_col=ycol, target_size=(img_size, img_size), class_mode='categorical', color_mode='rgb')
    msg='{0:70s} for valid generator'.format(' ')
    print(msg, '\r', end='') # prints over on the same line
    valid_gen=gen.flow_from_dataframe(valid_df, x_col='filepaths', y_col=ycol, target_size=(img_size, img_size), class_mode='categorical', color_mode='rgb')
    # for the test_gen we want to calculate the batch size and test steps such that
    # this insures that we go through all the sample in the test set exactly once
    length=len(test_df)
    test_batch_size=sorted([int(length/n) for n in range(1,length+1) if length % n == 0])
    test_steps=int(length/test_batch_size)
    msg='{0:70s} for test generator'.format(' ')
    print(msg, '\r', end='') # prints over on the same line
    test_gen=gen.flow_from_dataframe(test_df, x_col='filepaths', y_col=ycol, target_size=(img_size, img_size), class_mode='categorical', color_mode='rgb')
    # from the generator we can get information we will need later
    classes=list(train_gen.class_indices.keys())
    class_indices=list(train_gen.class_indices.values())
    class_count=len(classes)
    labels=test_gen.labels
    return train_gen, test_gen, valid_gen, test_steps, class_count
```

Create a function to show example training images

- gen is the ImageDataGenerator containing the images to be displayed

```
In [11]: def show_image_samples(gen, modnum ):
    msg='Below are some example training images'
    print_in_color(msg)
    t_dict=gen.class_indices
    classes=list(t_dict.keys())
    images,labels=next(gen) # get a sample batch from the generator
    plt.figure(figsize=(25, 25))
    length=len(labels)
    if length<25:    #show maximum of 25 images
        r=length
    else:
        r=25
    for i in range(r):
        plt.subplot(5, 5, i + 1)
        if modnum == 6:
            image=(images[i]+1)/2
        else:
            image=images[i] /255
        plt.imshow(image)
        index=np.argmax(labels[i])
        class_name=classes[index]
```

```
    plt.title(class_name, color='blue', fontsize=18)
    plt.axis('off')
plt.show()
```

Create a function to calculate the F1 score metric

- y_true is an np array containing the true integer index of the associated image file
- y_pred is an np array containing the predicted integer index of a test image in the test dataframe

```
In [12]: def F1_score(y_true, y_pred): #taken from old keras source code
    true_positives = K.sum(K.round(K.clip(y_true * y_pred, 0, 1)))
    possible_positives = K.sum(K.round(K.clip(y_true, 0, 1)))
    predicted_positives = K.sum(K.round(K.clip(y_pred, 0, 1)))
    precision = true_positives / (predicted_positives + K.epsilon())
    recall = true_positives / (possible_positives + K.epsilon())
    f1_val = 2*(precision*recall)/(precision+recall+K.epsilon())
    return f1_val
```

Create a model using transfer learning

- img_size is a tuple (height, width)
- class_count is an integer that specifies the number of classes in the dataset
- lr is a float specifying the intial model's learning rate
- ans is a string. If ans is an empty string an EfficientNetB0 model is used for transfer learning,
is ans='S' a MobilenetV3-small model is used for transfer learning. If ans='L'
an EfficientNetB3 model is used for transfer learning

```
In [13]: def make_model(img_size,class_count, lr, modnum):
    img_shape=(img_size[0], img_size[1], 3)
    if modnum ==1:
        base_model=tf.keras.applications.MobileNetV3Small(include_top=False, wei
        msg= 'created MobileNet V3 small model'
    elif modnum == 2:
        base_model=tf.keras.applications.MobileNetV3Large(include_top=False, wei
        msg='created MobileNetV3 large model'
    elif modnum == 3:
        base_model=tf.keras.applications.EfficientNetV2B0(include_top=False, wei
        msg='Created EfficientNetV2 B0 model'
    elif modnum == 4:
        base_model=tf.keras.applications.EfficientNetV2B1(include_top=False, wei
        msg='Created EfficientNetV2 B1 model'
    elif modnum == 5:
        base_model=tf.keras.applications.EfficientNetV2B2(include_top=False, wei
        msg='Created EfficientNetV2 B2 model'
    elif modnum == 6:
        tf.keras.applications.inception_resnet_v2.preprocess_input
```

```

base_model=tf.keras.applications.InceptionResNetV2(include_top=False, we
    msg='Created InceptionResNetV2 model'
base_model.trainable=True
x=base_model.output
x=BatchNormalization(axis=-1, momentum=0.99, epsilon=0.001 )(x)
x = Dense(256, kernel_regularizer = regularizers.l2(l = 0.016),activity_regu
                bias_regularizer=regularizers.l1(0.006) ,activation='relu')(x)
x=Dropout(rate=.4, seed=123)(x)
output=Dense(class_count, activation='softmax')(x)
model=Model(inputs=base_model.input, outputs=output)
model.compile(Adamax(learning_rate=lr), loss='categorical_crossentropy', met
msg=msg + f' with initial learning rate set to {lr}'
print_in_color(msg)
return model

```

Create a custom Keras callback to continue or halt training

The LR_ASK callback is a convenient callback that allows you to continue training for ask_epoch more epochs or to halt training.

If you elect to continue training for more epochs you are given the option to retain the current learning rate (LR) or to enter a new value for the learning rate. The form of use is:
ask=LR_ASK(model,epochs, ask_epoch) where:

- model is a string which is the name of your compiled model
- epochs is an integer which is the number of epochs to run specified in model.fit
- ask_epoch is an integer. If ask_epoch is set to a value say 5 then the model will train for 5 epochs.
then the user is asked to enter H to halt training, or enter an integer value. For example if you enter 4
training will continue for 4 more epochs to epoch 9 then you will be queried again. Once you enter an integer value you are prompted to press ENTER to continue training using the current learning rate
or to enter a new value for the learning rate.
- dwell is a boolean. If set to true the function compares the validation loss for the current to the lowest validation loss thus far achieved. If the validation loss for the current epoch is larger than learning rate
is automatically adjusted by the formula new_lr=lr * factor where factor is a float between 0 and 1. The motivation here is that if the validation loss increased we have moved to a point in Nspace on the cost function surface that is less favorable(higher cost) than for the epoch with the lowest cost. So the

model is loaded with the weights from the epoch with the lowest loss and the learning rate is reduced

At the end of training the model weights are set to the weights for the epoch that achieved the lowest validation loss

```
In [14]: class LR_ASK(keras.callbacks.Callback):
    def __init__(self, model, epochs, ask_epoch, batches, dwell=True, factor=.
        super(LR_ASK, self).__init__()
        self.model=model
        self.ask_epoch=ask_epoch
        self.epochs=epochs
        self.ask=True # if True query the user on a specified epoch
        self.lowest_vloss=np.inf
        self.lowest_aloss=np.inf
        self.best_weights=self.model.get_weights() # set best weights to model's
        self.best_epoch=1
        self.dwell=dwell
        self.factor=factor
        self.header=True
        self.batches=batches

    def on_train_begin(self, logs=None): # this runs on the beginning of trainin
        msg1 = f'Training will proceed until epoch {self.ask_epoch} then you will
        msg2='enter H to halt training or enter an integer for how many more epo
        print_in_color(msg1 + msg2)
        if self.dwell:
            msg='learning rate will be automatically adjusted during training'
            print_in_color(msg, (0,255,0))
            self.start_time= time.time() # set the time at which training started

    def on_train_end(self, logs=None): # runs at the end of training
        msg=f'loading model with weights from epoch {self.best_epoch}'
        print_in_color(msg, (0,255,255))
        self.model.set_weights(self.best_weights) # set the weights of the model
        tr_duration=time.time() - self.start_time # determine how long the tra
        hours = tr_duration // 3600
        minutes = (tr_duration - (hours * 3600)) // 60
        seconds = tr_duration - ((hours * 3600) + (minutes * 60))
        msg = f'training elapsed time was {str(hours)} hours, {minutes:.1f} min
        print_in_color (msg) # print out training duration time

    def on_epoch_begin(self, epoch, logs=None):
        self.ep_start = time.time()
    def on_train_batch_end(self, batch, logs=None):
        # get batch accuracy and loss
        acc = logs.get('accuracy') * 100
        loss = logs.get('loss')
        # prints over on the same line to show running batch count
        msg = '{0:20s}processing batch {1:} of {2:5s}- accuracy= {3:5.3f} - los
        print(msg, '\r', end='')

    def on_epoch_end(self, epoch, logs=None): # method runs on the end of each
        if self.header == True:
            msg = '{0:^7s}{1:^9s}{2:^9s}{3:^9s}{4:^10s}{5:^13s}{6:^10s}{7:^13s}{
            msg1=msg.format('Epoch', 'Train', 'Train', 'Valid', 'Valid', 'V_Loss
            msg='{0:^7s}{1:^9s}{2:^9s}{3:^9s}{4:^10s}{5:^13s}{6:^10s}{7:^13s}{8:
            msg2=msg.format(' ', 'Loss', 'Accuracy', 'Loss', 'Accuracy', 'Improve
```

```

        print_in_color (msg1 + msg2)
        self.header=False
    ep_end = time.time()
    duration = ep_end - self.ep_start
    vloss=logs.get('val_loss') # get the validation Loss for this epoch
    aloss=logs.get('loss')
    acc = logs.get('accuracy') # get training accuracy
    v_acc = logs.get('val_accuracy') # get validation accuracy
    lr = float(tf.keras.backend.get_value(self.model.optimizer.lr)) # get the
if epoch >0:
    deltv = self.lowest_vloss- vloss
    pimprov=(deltv/self.lowest_vloss) * 100
    deltaa=self.lowest_aloss-aloss
    aimprov=(deltaa/self.lowest_aloss) * 100
else:
    pimprov=0.0
if vloss< self.lowest_vloss:
    self.lowest_vloss=vloss
    self.best_weights=self.model.get_weights() # set best weights to model
    self.best_epoch=epoch + 1
    new_lr=lr
    msg = '{0:^7s}{1:^9.4f}{2:^9.2f}{3:^9.4f}{4:^10.2f}{5:^13.2f}{6:^10.2f}'.format(str(epoch+1), aloss, acc*100, vloss, v_acc*100, pimprov)
    print_in_color(msg, (0,255,0)) # green foreground
else: # validation loss increased
    if self.dwell: # if dwell is True when the validation loss increases
        lr=float(tf.keras.backend.get_value(self.model.optimizer.lr)) #
        new_lr=lr * self.factor
        msg = '{0:^7s}{1:^9.4f}{2:^9.2f}{3:^9.4f}{4:^10.2f}{5:^13.2f}{6:^10.2f}'.format(str(epoch+1), aloss, acc*100, vloss, v_acc*100, pimprov)
        print_in_color(msg, (255,255,0))
        tf.keras.backend.set_value(self.model.optimizer.lr, new_lr) # set the learning rate
        self.model.set_weights(self.best_weights) # set the weights of the model
if self.ask: # are the conditions right to query the user?
    if epoch + 1 ==self.ask_epoch: # is this epoch the one for querying the user?
        msg='\n Enter H to end training or an integer for the number of epochs to train'
        print_in_color(msg) # cyan foreground
        ans=input()
        if ans == 'H' or ans =='h' or ans == '0': # quit training for this epoch
            msg=f'you entered {ans}, Training halted on epoch {epoch+1}'
            print_in_color(msg)
            self.model.stop_training = True # halt training
        else: # user wants to continue training
            self.header=True
            self.ask_epoch += int(ans)
            msg=f'you entered {ans} Training will continue to epoch {self.ask_epoch}'
            print_in_color(msg) # cyan foreground
            if self.dwell==False:
                lr=float(tf.keras.backend.get_value(self.model.optimizer.lr))
                msg=f'current LR is {lr:8.6f} hit enter to keep this learning rate'
                print_in_color(msg) # cyan foreground
                ans=input(' ')
                if ans =='':
                    msg=f'keeping current LR of {lr:7.5f}'
                    print_in_color(msg) # cyan foreground
                else:
                    new_lr=float(ans)
                    tf.keras.backend.set_value(self.model.optimizer.lr, new_lr)
                    msg=f' changing LR to {ans}'
                    print_in_color(msg) # cyan foreground

```

Define a function to plot the training data

- tr_data is the history data from model.fit

```
In [15]: def tr_plot(tr_data):
    start_epoch=0
    #Plot the training and validation data
    tacc=tr_data.history['accuracy']
    tloss=tr_data.history['loss']
    vacc=tr_data.history['val_accuracy']
    vloss=tr_data.history['val_loss']
    tf1=tr_data.history['F1_score']
    vf1=tr_data.history['val_F1_score']
    Epoch_count=len(tacc)+ start_epoch
    Epochs=[]
    for i in range (start_epoch ,Epoch_count):
        Epochs.append(i+1)
    index_loss=np.argmin(vloss)# this is the epoch with the lowest validation loss
    val_lowest=vloss[index_loss]
    index_acc=np.argmax(vacc)
    acc_highest=vacc[index_acc]
    indexf1=np.argmax(vf1)
    vf1_highest=vf1[indexf1]
    plt.style.use('fivethirtyeight')
    sc_label='best epoch= '+ str(index_loss+1 +start_epoch)
    vc_label='best epoch= '+ str(index_acc + 1+ start_epoch)
    f1_label='best epoch= '+ str(index_acc + 1+ start_epoch)
    fig,axes=plt.subplots(nrows=1, ncols=3, figsize=(25,10))
    axes[0].plot(Epochs,tloss, 'r', label='Training loss')
    axes[0].plot(Epochs,vloss,'g',label='Validation loss' )
    axes[0].scatter(index_loss+1 +start_epoch,val_lowest, s=150, c= 'blue', label=sc_label)
    axes[0].scatter(Epochs, tloss, s=100, c='red')
    axes[0].set_title('Training and Validation Loss')
    axes[0].set_xlabel('Epochs', fontsize=18)
    axes[0].set_ylabel('Loss', fontsize=18)
    axes[0].legend()
    axes[1].plot (Epochs,tacc,'r',label= 'Training Accuracy')
    axes[1].scatter(Epochs, tacc, s=100, c='red')
    axes[1].plot (Epochs,vacc,'g',label= 'Validation Accuracy')
    axes[1].scatter(index_acc+1 +start_epoch,acc_highest, s=150, c= 'blue', label=vc_label)
    axes[1].set_title('Training and Validation Accuracy')
    axes[1].set_xlabel('Epochs', fontsize=18)
    axes[1].set_ylabel('Accuracy', fontsize=18)
    axes[1].legend()
    axes[2].plot (Epochs,tf1,'r',label= 'Training F1 score')
    axes[2].plot (Epochs,vf1,'g',label= 'Validation F1 score')
    index_tf1=np.argmax(tf1)# this is the epoch with the highest training F1 score
    tf1max=tf1[index_tf1]
    index_vf1=np.argmax(vf1)# this is the epoch with the highest validation F1 score
    vf1max=vf1[index_vf1]
    axes[2].scatter(index_vf1+1 +start_epoch,vf1max, s=150, c= 'blue', label=vc_label)
    axes[2].scatter(Epochs, tf1, s=100, c='red')
    axes[2].set_title('Training and Validation F1 score')
    axes[2].set_xlabel('Epochs', fontsize=18)
```

```

    axes[2].set_ylabel('F1 score', fontsize=18)
    axes[2].legend()
    plt.tight_layout()
    plt.show()
    return

```

Define a function to save the training data to a csv file

- history is the history from history=model.fit
- csvpath is a string specifying the full path where the csv file will be stored

```
In [16]: def save_history_to_csv(history,csvpath):
    trdict=history.history
    df=pd.DataFrame()
    df['Epoch']=list(np.arange(1, len(trdict['loss']) + 1 ))
    keys=list(trdict.keys())
    for key in keys:
        data=list(trdict[key])
        df[key]=data
    df.to_csv(csvpath, index=False)
```

Make Predictions on the test set

Define a function which takes in a trained model and a test generator and generates predictions

on the test set including a confusion matrix and a classification report

- model is the trained model
- test_gen is the ImageDataGenerator holding the test filepaths and labels

The function will produce a classification report, a confusion matrix and a plot of misclassifications

```
In [17]: def predictor(model,test_gen):
    classes=list(test_gen.class_indices.keys())
    class_count=len(classes)
    preds=model.predict(test_gen, verbose=1)
    errors=0
    test_count =len(preds)
    misclassified_classes=[]
    misclassified_files=[]
    misclassified_as = []
    pred_indices=[]
    for i, p in enumerate (preds):
        pred_index=np.argmax(p)
        pred_indices.append(pred_index)
        true_index= test_gen.labels[i]
```

```

if pred_index != true_index:
    errors +=1
    misclassified_classes.append(classes[true_index])
    misclassified_as.append(classes[pred_index])
    file=test_gen.filenames[i]
    split=file.split('/')
    L=len(split)
    f=split[L-2] +' '+ split[L-1]
    misclassified_files.append(f)

accuracy = (test_count-errors)*100/test_count
ytrue=np.array(test_gen.labels)
ypred=np.array(pred_indices)
f1score=f1_score(ytrue, ypred, average='weighted')* 100
msg=f'There were {errors} errors in {test_count} tests for an accuracy of {a}
print (msg)
misclassified_classes=sorted(misclassified_classes)
if len(misclassified_classes) > 0:
    misclassifications=[]
    for klass in misclassified_classes:
        mis_count=misclassified_classes.count(klass)
        misclassifications.append(mis_count)
    unique=len(np.unique(misclassified_classes))
    if unique==1:
        height=int(unique)
    else:
        height =int(unique/2)
    plt.figure(figsize=(10, height))
    plt.style.use('fivethirtyeight')
    plt.barh(misclassified_classes, misclassifications )
    plt.title( 'Classification Errors on Test Set by Class', fontsize=20, co
    plt.xlabel('NUMBER OF MISCLASSIFICATIONS', fontsize=20, color='blue')
    plt.ylabel('CLASS', fontsize=20, color='blue')
    plt.show()
if class_count <=30:
    cm = confusion_matrix(ytrue, ypred )
    # plot the confusion matrix
    plt.figure(figsize=(12, 8))
    sns.heatmap(cm, annot=True, vmin=0, fmt='g', cmap='Blues', cbar=False)
    plt.xticks(np.arange(class_count)+.5, classes, rotation=90)
    plt.yticks(np.arange(class_count)+.5, classes, rotation=0)
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.title("Confusion Matrix")
    plt.show()
    clr = classification_report(ytrue, ypred, target_names=classes, digits=
    print("Classification Report:\n-----\n", clr)
return f1score, misclassified_files

```

Save the model

In [18]:

```

def save_model(model,subject, classes, img_size, f1score, working_dir):
    name=f'{subject}-{str(len(classes))}-{str(img_size[0])} X {str(img_size[1])}'
    model_save_loc=os.path.join(working_dir, name)
    try:
        model.save(model_save_loc)
        msg= f'model was saved as {model_save_loc}'
    
```

```

        print_in_color(msg, (0,255,255),(0,0,0))
except:
    msg='model can not be saved due to tensorflow 2.10.0 or higher. Bug invo
    print_in_color(msg, (0,255,255), (0,0,0)) # cyan foreground

```

Define a function that runs the classifier

This is the code that provides the interface to the user and calls the various defined functions based on the users input. To follow how the program works start here. You can run this function as many times as you care to if you want to do runs with various different parameters. You can save the data from each run in a csv file to compare training and validation results from different runs.

```

In [19]: def run(working_dir= None,train_dir=None, test_dir=None, valid_dir=None,start=True):
    if start:
        msg='enter the full path to the working directory where data will be sto
        print_in_color(msg)
        working_dir=input(' ')
        if 'kaggle' in working_dir: # if running the notebook on kaggle need to
            delimiter='/'
        else:
            delimiter='\\'
        msg='Enter the full path to the train directory'
        print_in_color(msg)
        train_dir =input(' ')
        msg = 'Enter the full path to the validation directory. If there is no v
        print_in_color(msg)
        ans=input(' ')
        if ans == '':
            valid_dir=None
        else:
            valid_dir=ans
        msg = 'Enter the full path to the test directory. If there is no test di
        print_in_color(msg)
        ans=input(' ')
        if ans == '':
            test_dir=None
        else:
            test_dir=ans

    ftotal, flargest, maxclass , fsmallest, minclass= check_dataset_size(train_d
    msg1= f'the train directory contains {ftotal} files, class {maxclass} has th
    msg2= f'class {minclass} contains the least number of images of {fsmallest}
    msg3= f'NOTE if the value of most images is the same as the value of least i
    msg4='When dealing with very large data sets to save time you may not want t
    msg5 = 'to limit the maximum number of image files in any class you can enter
    msg=msg1 +msg2 + msg3 + msg4 + msg5
    print_in_color(msg) # print the training data information for the user
    msg = 'input a limiter integer value to limit max number of images in a clas
    print_in_color(msg)
    limiter=input(' ')
    if limiter == '':
        limiter = None

```

```

        msg='dataset will be processed as is with no limiter'
        print_in_color(msg)
else:
    limiter=int(limiter)
    msg=f'images will be limited to a maximum of {limiter} images in each cl
    print_in_color(msg)
msg='enter the minimum number of training images a class must have to be inc
print_in_color(msg)
min_images=int(input(' '))
if min_images=='':
    min_images=10
else:
    min_images=int(min_images)
# create train, test and valid data frames
train_df, test_df, valid_df, classes, class_count, max_samples, min_samples,
msg=f'the average height of training images is {have} and average width is {
print_in_color(msg)
img_height=int(input(' '))
msg=' Enter the image width to be used to train the model'
print_in_color(msg)
img_width=int(input(' '))
img_size=(img_height, img_width)
msg=f'model will be trained with image shape of  ( {img_height}, {img_width})
print_in_color(msg)
# determine if user wants to balance the training dataframe
msg=' enter A  to auto balance the train set or enter or hit enter to leave
print_in_color(msg)
ans=input(' ')
if ans == 'A' or ans =='a':
    msg='enter the number of images you want to have in each class of the tr
    print_in_color(msg)
    max_images=int(input(' '))
    train_df, classes, class_count=trim (train_df, max_images, min_images, '
    train_df=balance(train_df, max_images, 'labels', working_dir, img_size)
    plot_title='Images per Label after Auto Balance of train data set'
    plot_label_count (train_df, plot_title) # plot the number of images in
else:
    msg=f'training data set will be used as is '
    print_in_color (msg)
classes=list(train_df['labels'].unique())
class_count = len(classes)
msg='enter the desired batch size or press enter to use the default batch si
print_in_color(msg)
ans=input(' ')
if ans == '':
    bs=20
else:
    bs=int(ans)
# make the train, test and valid data generators
print ('select a model number from the list below ',flush=True)
print ('{:^9s}{:^20s}{:^12s}'.format('Model No', 'Model Type', 'Parameter
models=['MobileNetV3-small', 'MobileNetV3-large', 'EfficientNetV2B0', 'Effic
parameters=[' 1.0 M', '3.3 M', '6.3 M', '7.3 M','9.1 M', '54.7 M']
for i in range(0,6):
    msg='{:^9s}{:^20s}{:^12s}'.format(str(i+1),models[i], parameters[i])
    print (msg)
modnum=int(input(" "))
lr=.001
model=make_model(img_size,class_count, lr, modnum)
msg1='enter the number of epochs to run initially. After these epochs comple

```

```

msg2='or enter an integer for how many more epochs to run then be asked again'
print(msg1 + msg2)
ask_epoch=int(input(' '))
batches=int(len(train_df)/bs)
# instantiate the custom callback
epochs=100
ask=LR_ASK(model, epochs=epochs, ask_epoch=ask_epoch, batches=batches) # in
callbacks=[ask]
# make the train, test and valid data generators
train_gen, test_gen, valid_gen, test_steps, class_count= make_gens(bs, 'labe
show_image_samples(train_gen, modnum ) # show some sample training images
# train the model
history=model.fit(x=train_gen, epochs=epochs, verbose=0, callbacks=callback
validation_steps=None, shuffle=True, initial_epoch=0) # tra

tr_plot(history) # plot training data
msg='To save the training data to a csv file enter the name for the csv file
print_in_color(msg)
ans=input(' ')
if ans !='': # save the training data to a csv file history.csv in working d
    csvpath=os.path.join(working_dir, ans + '.csv')
    save_history_to_csv(history,csvpath)
    msg=f'training data saved to {csvpath}'
    print_in_color(msg)
# make predictions on the test set, create classification report and confusi
f1score, misclassified_files=predictor(model,test_gen) # do predictions on t
if len(misclassified_files)>0:
    print ('below is a list of misclassified test files:')
    for f in misclassified_files:
        print (f)
# save the model
msg=f'your trained model will be saved to directory {working_dir} enter a su
print_in_color(msg)
subject=input(' ')
save_model(model,subject, classes, img_size, f1score, working_dir) # save th
msg='model save nomenclature is directory/subject-number of classes- (img_he
print_in_color(msg)
# remove the aug directory from the working directory
if os.path.isdir(os.path.join(working_dir, 'aug')):
    shutil.rmtree(os.path.join(working_dir, 'aug')) #remove the augmentation
return working_dir,train_dir, test_dir, valid_dir, model

```

Initiate the run function and enable ability to execute run multiple times

The code below initiate the first run of the classifier and lets the user rerun the classifier

Save the model

```

In [20]: working_dir,train_dir, test_dir, valid_dir, model=run(start=True) # run the clas
stop = False
while stop == False:
    msg='enter R to rerun the classifier or press enter to quit '

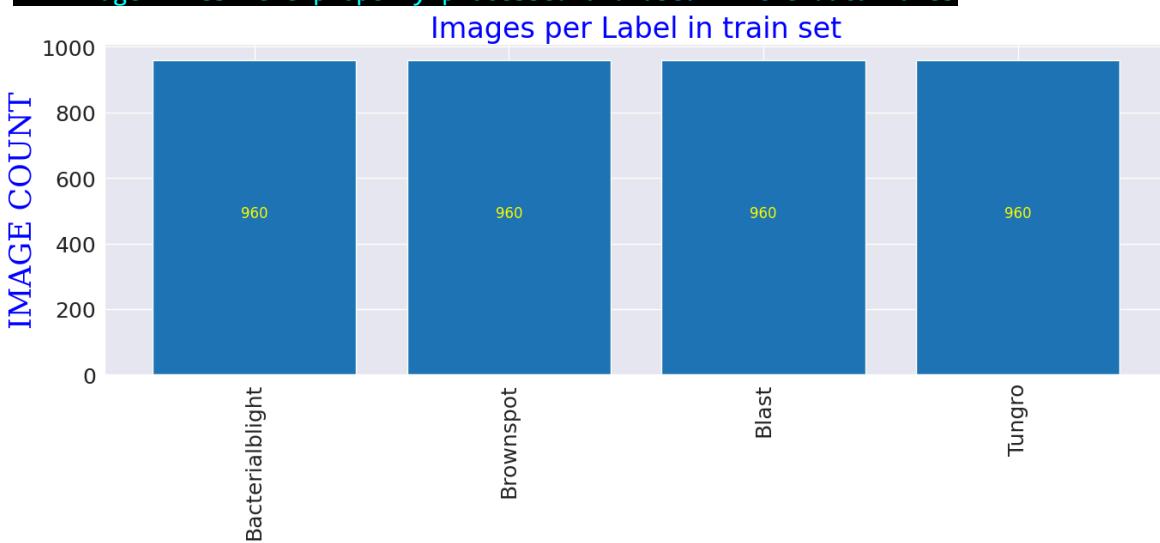
```

```

print_in_color(msg)
ans=input(' ')
if ans == 'R' or ans == 'r':
    msg='Enter N to specify new data directories or hit Enter to use existin
    print_in_color(msg)
    ans=input(' ')
    if ans == '':
        run(working_dir,train_dir=train_dir, test_dir=test_dir, valid_dir=va
    else:
        run(True)
else:
    stop = True
    print_in_color ('process commpleted')

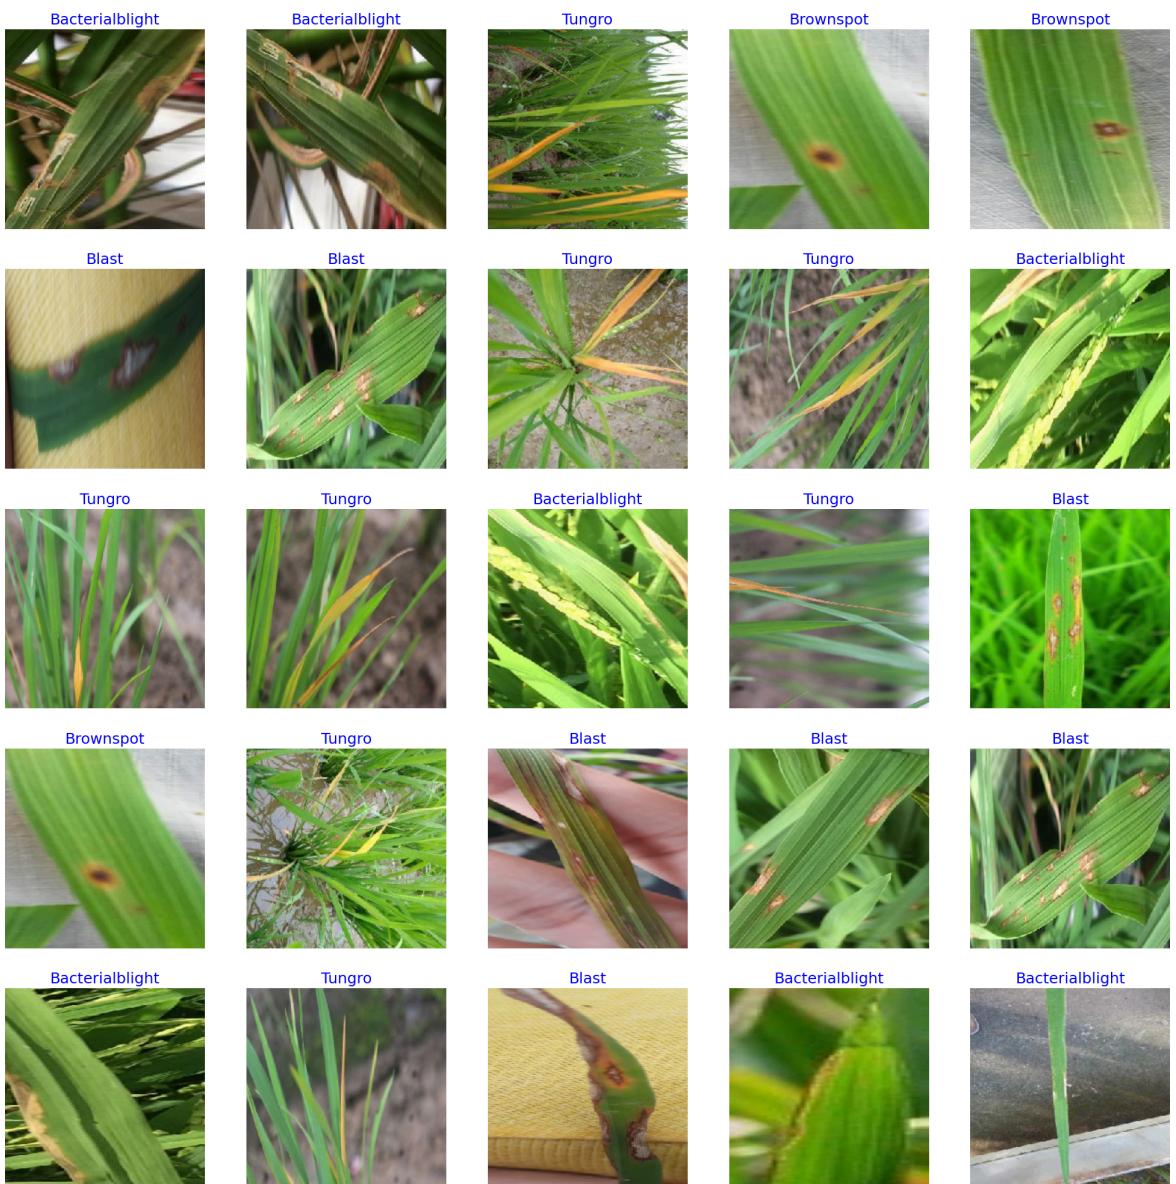
```

enter the full path to the working directory where data will be stored.
 Enter the full path to the train directory
 Enter the full path to the validation directory. If there is no validation direct
 ory press enter
 Enter the full path to the test directory. If there is no test directory press en
 ter
 the train directory contains 5932 files, class Brownspot has the most images of 1
 600 files
 class Tungro contains the least number of images of 1308 files
 NOTE if the value of most images is the same as the value of least images the dat
 asset is balanced
 When dealing with very large data sets to save time you may not want to read in a
 ll the image files
 to limit the maximum number of image files in any class you can enter an integer
 limiter value
 input a limiter integer value to limit max number of images in a class, for no li
 mit hit enter
 images will be limited to a maximum of 1200 images in each class
 enter the minimum number of training images a class must have to be included in t
 he dataset or press enter to use the default value of 10
 processing images in train directory for class Tungro
 number of classes in processed dataset= 4
 the maximum files in any class in train_df is 960 the minimum files in any cla
 ss in train_df is 960
 train_df length: 3840 test_df length: 480 valid_df length: 480
 All image files were properly processed and used in the dataframes



the average height of training images is 313 and average width is 314
 enter the image height to be used to train the model
 Enter the image width to be used to train the model

```
model will be trained with image shape of ( 224, 224 )
enter A to auto balance the train set or enter or hit enter to leave train set
unchanged
training data set will be used as is
enter the desired batch size or press enter to use the default batch size of 20
select a model number from the list below
Model No      Model Type      Parameters
1      MobileNetV3-small    1.0 M
2      MobileNetV3-large    3.3 M
3      EfficientNetV2B0     6.3 M
4      EfficientNetV2B1     7.3 M
5      EfficientNetV2B2     9.1 M
6      InceptionResNetV2   54.7 M
Downloading data from https://storage.googleapis.com/tensorflow/keras-application-
s/efficientnet_v2/efficientnetv2-b0_notop.h5
24274472/24274472 [=====] - 0s 0us/step
Created EfficientNetV2 B0 model with initial learning rate set to 0.001
enter the number of epochs to run initially. After these epochs complete you can
enter H to halt training
or enter an integer for how many more epochs to run then be asked again
Found 3840 validated image filenames belonging to 4 classes.          for train
generator
Found 480 validated image filenames belonging to 4 classes.          for valid
generator
Found 480 validated image filenames belonging to 4 classes.          for test g
enerator
Below are some example training images
```



**Training will proceed until epoch 10 then you will be asked to
enter H to halt training or enter an integer for how many more epochs to run then
be asked again
learning rate will be automatically adjusted during training**

2024-01-01 12:03:41.314932: E tensorflow/core/grappler/optimizers/meta_optimizer.
cc:954] layout failed: INVALID_ARGUMENT: Size of values 0 does not match size of
permutation 4 @ fanin shape inmodel/block2b_drop/dropout/SelectV2-2-TransposeNHWC
ToNCHW-LayoutOptimizer

Epoch	Train Duration in Seconds	Train Loss	Train Accuracy	Valid Loss	Valid Accuracy	V_Loss %	Learning Rate	Next LR	D
1	6.6278	90.47	5.2828	99.58	0.00	0.001000	0.001000		
2	23.55	4.3160	99.11	3.5387	100.00	33.01	0.001000	0.001000	
3	23.75	2.9557	99.69	2.4282	100.00	31.38	0.001000	0.001000	
4	23.50	2.0656	99.87	1.6991	100.00	30.03	0.001000	0.001000	
5	23.33	1.4564	99.95	1.1903	100.00	29.94	0.001000	0.001000	
6	23.34	1.0359	99.97	0.8439	100.00	29.10	0.001000	0.001000	
7	23.41	0.7431	99.95	0.6039	100.00	28.44	0.001000	0.001000	
8	23.54	0.5456	99.92	0.4427	100.00	26.70	0.001000	0.001000	
9	23.46	0.4057	100.00	0.3285	100.00	25.80	0.001000	0.001000	
10	23.39	0.3111	99.95	0.2542	100.00	22.61	0.001000	0.001000	

Enter H to end training or an integer for the number of additional epochs to run then ask again
you entered h, Training halted on epoch 10 due to user input

loading model with weights from epoch 10
training elapsed time was 0.0 hours, 4.0 minutes, 53.01 seconds)



To save the training data to a csv file enter the name for the csv file or press enter to not save the data

6/6 [=====] - 4s 158ms/step

There were 0 errors in 480 tests for an accuracy of 100.00 and an F1 score of 10.00

		Confusion Matrix			
		Actual		Predicted	
		Bacterialblight	Blast	Brownspot	Tungro
Bacterialblight	120	0	0	0	0
Blast	0	120	0	0	0
Brownspot	0	0	120	0	0
Tungro	0	0	0	120	

Classification Report:

	precision	recall	f1-score	support
Bacterialblight	1.0000	1.0000	1.0000	120
Blast	1.0000	1.0000	1.0000	120
Brownspot	1.0000	1.0000	1.0000	120
Tungro	1.0000	1.0000	1.0000	120
accuracy			1.0000	480
macro avg	1.0000	1.0000	1.0000	480
weighted avg	1.0000	1.0000	1.0000	480

```

your trained model will be saved to directory /kaggle/working/ enter a subject for the saved model
model was saved as /kaggle/working/EfficientNetB0-4-(224 X 224)- 100.00.h5
model save nomenclature is directory/subject-number of classes- (img_height, img_width)- F1score.h5
enter R to rerun the classifier or press enter to quit
Enter N to specify new data directories or hit Enter to use existing data directory paths
the train directory contains 5932 files, class Brownspot has the most images of 1600 files
class Tungro contains the least number of images of 1308 files
NOTE if the value of most images is the same as the value of least images the dataset is balanced
When dealing with very large data sets to save time you may not want to read in all the image files
to limit the maximum number of image files in any class you can enter an integer limiter value
input a limiter integer value to limit max number of images in a class, for no limit hit enter
images will be limited to a maximum of 1200 images in each class
enter the minimum number of training images a class must have to be included in the dataset or press enter to use the default value of 10

```

processing images in train directory for class Tungro
 number of classes in processed dataset= 4
 the maximum files in any class in train_df is 960 the minimum files in any class in train_df is 960
 train_df length: 3840 test_df length: 480 valid_df length: 480
All image files were properly processed and used in the dataframes



the average height of training images is 313 and average width is 314
enter the image height to be used to train the model
Enter the image width to be used to train the model
model will be trained with image shape of (224, 224)
enter A to auto balance the train set or enter or hit enter to leave train set unchanged
training data set will be used as is
enter the desired batch size or press enter to use the default batch size of 20
 select a model number from the list below

Model No	Model Type	Parameters
1	MobileNetV3-small	1.0 M
2	MobileNetV3-large	3.3 M
3	EfficientNetV2B0	6.3 M
4	EfficientNetV2B1	7.3 M
5	EfficientNetV2B2	9.1 M
6	InceptionResNetV2	54.7 M

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v3/weights_mobilenet_v3_small_224_1.0_float_no_top_v2.h5
 4334752/4334752 [=====] - 0s 0us/step
created MobileNet V3 small model with initial learning rate set to 0.001
 enter the number of epochs to run initially. After these epochs complete you can enter H to halt training
 or enter an integer for how many more epochs to run then be asked again
 Found 3840 validated image filenames belonging to 4 classes. for train generator
 Found 480 validated image filenames belonging to 4 classes. for valid generator
 Found 480 validated image filenames belonging to 4 classes. for test generator
Below are some example training images



Training will proceed until epoch 10 then you will be asked to
enter H to halt training or enter an integer for how many more epochs to run then
be asked again

learning rate will be automatically adjusted during training

Epoch	Train Duration in Seconds	Train Loss	Train Accuracy	Valid Loss	Valid Accuracy	V_Loss %	Learning Rate	Next LR	D
1	5.6640	88.67	4.8065	85.00	0.00	0.001000	0.001000		
37.67									
2	3.7287	99.32	3.5380	72.50	26.39	0.001000	0.001000		
11.16									
3	2.5715	99.74	2.6927	70.21	23.89	0.001000	0.001000		
10.57									
4	1.7984	99.97	2.0817	69.79	22.69	0.001000	0.001000		
11.80									
5	1.2703	99.95	1.5470	74.37	25.69	0.001000	0.001000		
10.82									
6	0.9047	99.95	1.2278	75.83	20.63	0.001000	0.001000		
10.34									
7	0.6511	99.97	0.9282	81.04	24.41	0.001000	0.001000		
10.71									
8	0.4811	100.00	0.7125	85.62	23.24	0.001000	0.001000		
10.45									
9	0.3627	100.00	0.5460	90.00	23.36	0.001000	0.001000		
10.50									
10	0.2814	99.97	0.4500	91.67	17.59	0.001000	0.001000		
10.22									

Enter H to end training or an integer for the number of additional epochs to run then ask again

you entered 5 Training will continue to epoch 15

Epoch	Train Duration in Seconds	Train Loss	Train Accuracy	Valid Loss	Valid Accuracy	V_Loss %	Learning Rate	Next LR	D
11	0.2290	99.97	0.3587	94.17	20.28	0.001000	0.001000		
10.52									
12	0.1913	100.00	0.2830	96.25	21.10	0.001000	0.001000		
11.33									
13	0.1665	99.97	0.2335	96.88	17.49	0.001000	0.001000		
11.90									
14	0.1518	99.97	0.1764	98.33	24.47	0.001000	0.001000		
10.41									
15	0.1358	100.00	0.1398	100.00	20.77	0.001000	0.001000		
10.92									

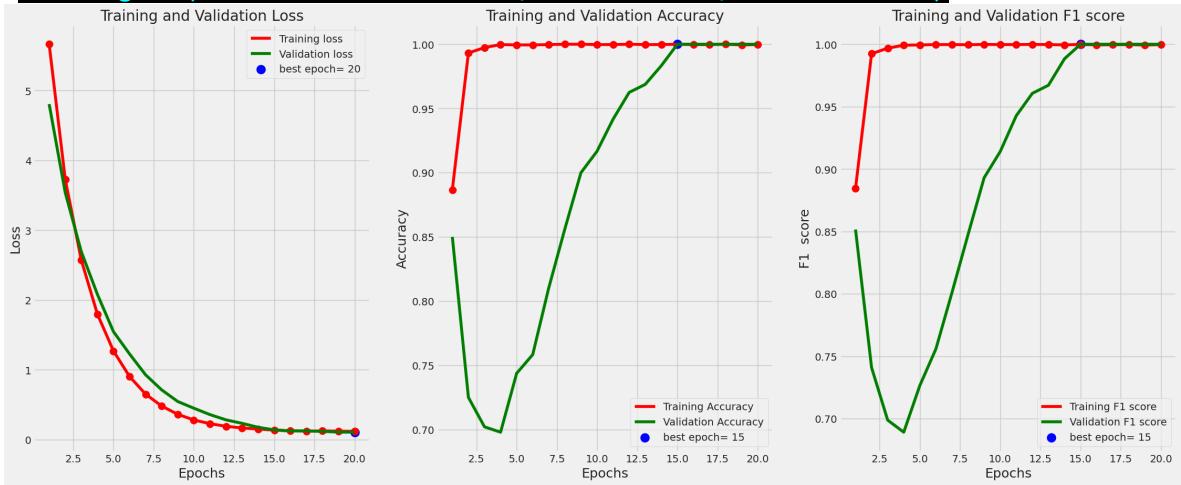
Enter H to end training or an integer for the number of additional epochs to run then ask again

you entered 5 Training will continue to epoch 20								
Epoch	Train duration in Seconds	Train Loss	Train Accuracy	Valid Loss	Valid Accuracy	V_Loss % Improvement	Learning Rate	Next LR D
16	11.40	0.1271	99.97	0.1242	100.00	11.15	0.001000	0.001000
17	10.89	0.1226	99.97	0.1252	100.00	-0.81	0.001000	0.000400
18	11.01	0.1246	100.00	0.1178	100.00	5.16	0.000400	0.000400
19	10.62	0.1213	99.95	0.1082	100.00	8.10	0.000400	0.000400
20	10.84	0.1185	99.97	0.1066	100.00	1.51	0.000400	0.000400

Enter H to end training or an integer for the number of additional epochs to run then ask again
you entered h, Training halted on epoch 20 due to user input

loading model with weights from epoch 20

training elapsed time was 0.0 hours, 4.0 minutes, 18.96 seconds)



To save the training data to a csv file enter the name for the csv file or press enter to not save the data

6/6 [=====] - 3s 157ms/step

There were 0 errors in 480 tests for an accuracy of 100.00 and an F1 score of 100.00

		Confusion Matrix			
		Actual		Predicted	
		Bacterialblight	Blast	Brownspot	Tungro
Bacterialblight	120	0	0	0	0
Blast	0	120	0	0	0
Brownspot	0	0	120	0	0
Tungro	0	0	0	0	120

Classification Report:

	precision	recall	f1-score	support
Bacterialblight	1.0000	1.0000	1.0000	120
Blast	1.0000	1.0000	1.0000	120
Brownspot	1.0000	1.0000	1.0000	120
Tungro	1.0000	1.0000	1.0000	120
accuracy			1.0000	480
macro avg	1.0000	1.0000	1.0000	480
weighted avg	1.0000	1.0000	1.0000	480

```

your trained model will be saved to directory /kaggle/working/ enter a subject for the saved model
model was saved as /kaggle/working/MobileNetV3_small-4-(224 X 224)- 100.00.h5
model save nomenclature is directory/subject-number of classes- (img_height, img_width)- F1score.h5
enter R to rerun the classifier or press enter to quit
Enter N to specify new data directories or hit Enter to use existing data directory paths
the train directory contains 5932 files, class Brownspot has the most images of 1600 files
class Tungro contains the least number of images of 1308 files
NOTE if the value of most images is the same as the value of least images the dataset is balanced
When dealing with very large data sets to save time you may not want to read in all the image files
to limit the maximum number of image files in any class you can enter an integer limiter value
input a limiter integer value to limit max number of images in a class, for no limit hit enter
images will be limited to a maximum of 1200 images in each class
enter the minimum number of training images a class must have to be included in the dataset or press enter to use the default value of 10

```

processing images in train directory for class Tungro
 number of classes in processed dataset= 4
 the maximum files in any class in train_df is 960 the minimum files in any class in train_df is 960
 train_df length: 3840 test_df length: 480 valid_df length: 480
All image files were properly processed and used in the dataframes



the average height of training images is 313 and average width is 314
enter the image height to be used to train the model
Enter the image width to be used to train the model
model will be trained with image shape of (224, 224)
enter A to auto balance the train set or enter or hit enter to leave train set unchanged
training data set will be used as is
enter the desired batch size or press enter to use the default batch size of 20
 select a model number from the list below

Model No	Model Type	Parameters
1	MobileNetV3-small	1.0 M
2	MobileNetV3-large	3.3 M
3	EfficientNetV2B0	6.3 M
4	EfficientNetV2B1	7.3 M
5	EfficientNetV2B2	9.1 M
6	InceptionResNetV2	54.7 M

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v3/weights_mobilenet_v3_large_224_1.0_float_no_top_v2.h5
 12683000/12683000 [=====] - 0s 0us/step
created MobileNetV3 large model with initial learning rate set to 0.001
 enter the number of epochs to run initially. After these epochs complete you can enter H to halt training
 or enter an integer for how many more epochs to run then be asked again
 Found 3840 validated image filenames belonging to 4 classes. for train generator
 Found 480 validated image filenames belonging to 4 classes. for valid generator
 Found 480 validated image filenames belonging to 4 classes. for test generator
Below are some example training images



Training will proceed until epoch 10 then you will be asked to
enter H to halt training or enter an integer for how many more epochs to run then
be asked again

learning rate will be automatically adjusted during training

Epoch	Train Duration in Seconds	Train Loss	Train Accuracy	Valid Loss	Valid Accuracy	V_Loss %	Learning Rate	Next LR	D
1	47.75	6.2302	92.63	5.8153	68.12	0.00	0.001000	0.001000	
2	15.58	4.0382	99.87	3.8419	75.42	33.93	0.001000	0.001000	
3	15.32	2.7367	99.82	2.7654	78.33	28.02	0.001000	0.001000	
4	15.29	1.8964	99.97	2.0067	82.08	27.44	0.001000	0.001000	
5	15.38	1.3412	99.97	1.4581	86.87	27.34	0.001000	0.001000	
6	15.23	0.9612	99.97	1.1328	86.67	22.31	0.001000	0.001000	
7	15.23	0.6997	100.00	0.8419	90.00	25.68	0.001000	0.001000	
8	15.32	0.5180	100.00	0.5998	94.17	28.76	0.001000	0.001000	
9	15.27	0.3899	100.00	0.4347	96.88	27.52	0.001000	0.001000	
10	15.22	0.3017	100.00	0.3442	97.29	20.81	0.001000	0.001000	

Enter H to end training or an integer for the number of additional epochs to run then ask again

you entered 5 Training will continue to epoch 15

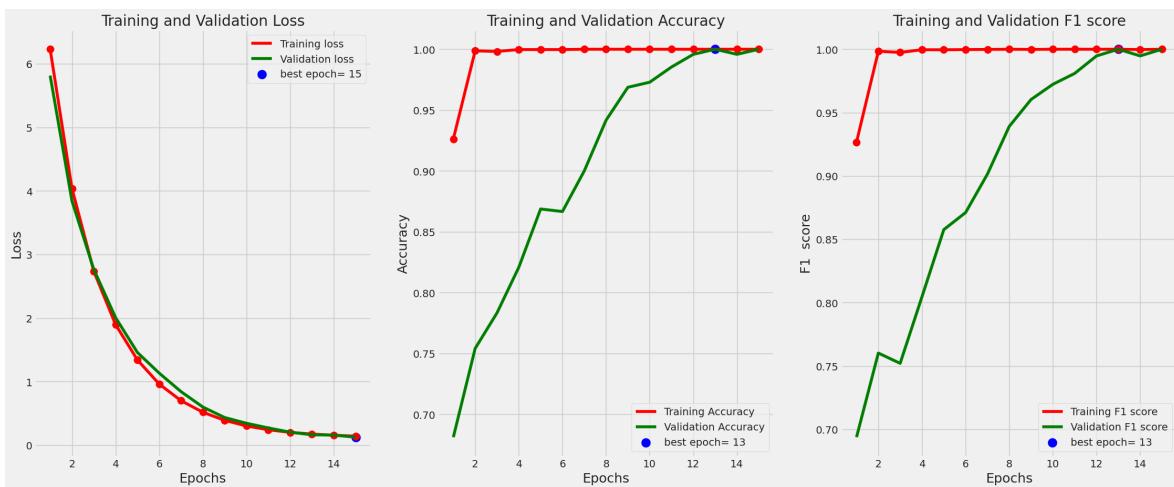
Epoch	Train Duration in Seconds	Train Loss	Train Accuracy	Valid Loss	Valid Accuracy	V_Loss %	Learning Rate	Next LR	D
11	15.25	0.2425	100.00	0.2707	98.54	21.37	0.001000	0.001000	
12	15.33	0.2010	100.00	0.2043	99.58	24.50	0.001000	0.001000	
13	15.45	0.1743	100.00	0.1629	100.00	20.26	0.001000	0.001000	
14	15.17	0.1561	100.00	0.1586	99.58	2.64	0.001000	0.001000	
15	15.38	0.1415	100.00	0.1253	100.00	21.02	0.001000	0.001000	

Enter H to end training or an integer for the number of additional epochs to run then ask again

you entered h, Training halted on epoch 15 due to user input

loading model with weights from epoch 15

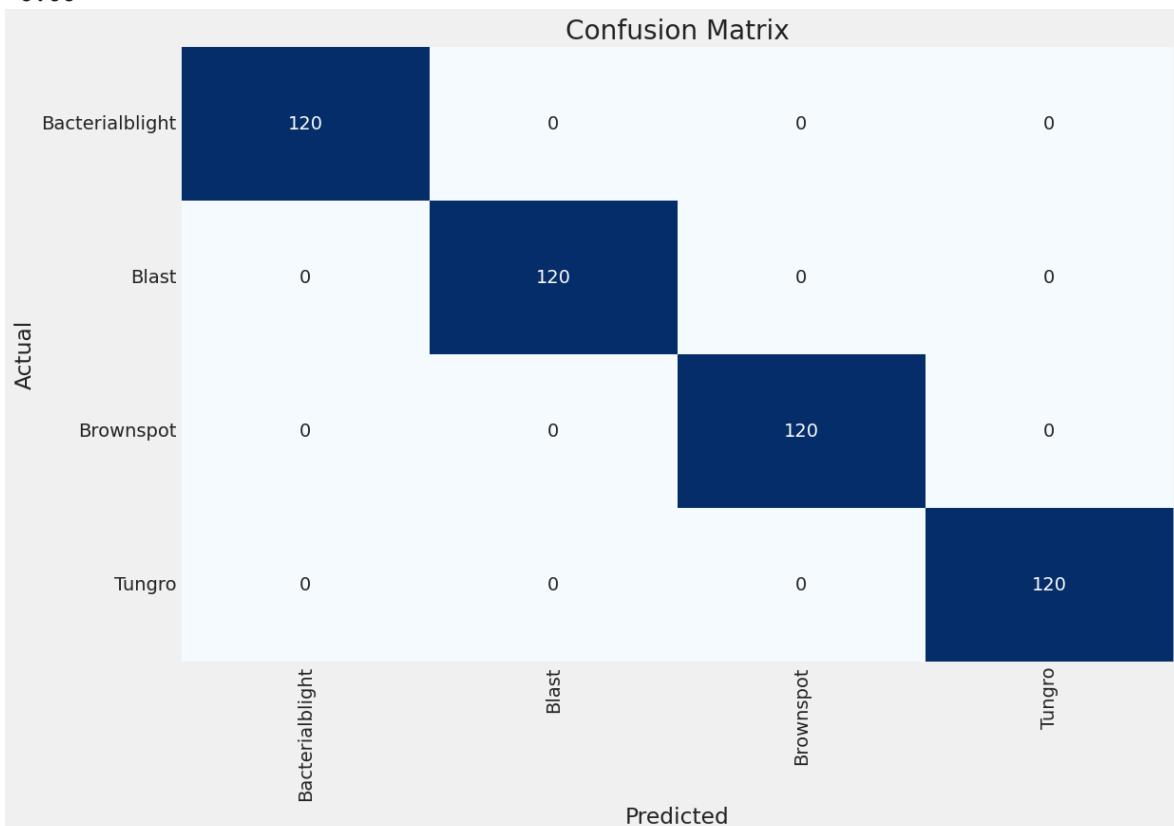
training elapsed time was 0.0 hours, 4.0 minutes, 34.04 seconds)



To save the training data to a csv file enter the name for the csv file or press enter to not save the data

6/6 [=====] - 3s 152ms/step

There were 0 errors in 480 tests for an accuracy of 100.00 and an F1 score of 100.00



Classification Report:

	precision	recall	f1-score	support
Bacterialblight	1.0000	1.0000	1.0000	120
Blast	1.0000	1.0000	1.0000	120
Brownspot	1.0000	1.0000	1.0000	120
Tungro	1.0000	1.0000	1.0000	120
accuracy			1.0000	480
macro avg	1.0000	1.0000	1.0000	480
weighted avg	1.0000	1.0000	1.0000	480

your trained model will be saved to directory /kaggle/working/ enter a subject for the saved model

```

model was saved as /kaggle/working/MobileNetV3 large-4-(224 X 224)- 100.00.h5
model save nomenclature is directory/subject-number of classes- (img_height, img_
width)- F1score.h5
enter R to rerun the classifier or press enter to quit
Enter N to specify new data directories or hit Enter to use existing data directo
ry paths
the train directory contains 5932 files, class Brownspot has the most images of 1
600 files
class Tungro contains the least number of images of 1308 files
NOTE if the value of most images is the same as the value of least images the dat
aset is balanced
When dealing with very large data sets to save time you may not want to read in a
ll the image files
to limit the maximum number of image files in any class you can enter an integer
limiter value
input a limiter integer value to limit max number of images in a class, for no li
mit hit enter
images will be limited to a maximum of 1200 images in each class
enter the minimum number of training images a class must have to be included in t
he dataset or press enter to use the default value of 10
processing images in train directory for class Tungro
number of classes in processed dataset= 4
the maximum files in any class in train_df is 960 the minimum files in any cla
ss in train_df is 960
train_df length: 3840 test_df length: 480 valid_df length: 480
All image files were properly processed and used in the dataframes

```



```

the average height of training images is 313 and average width is 314
enter the image height to be used to train the model
Enter the image width to be used to train the model
model will be trained with image shape of ( 224, 224 )
enter A to auto balance the train set or enter or hit enter to leave train set
unchanged
training data set will be used as is
enter the desired batch size or press enter to use the default batch size of 20
select a model number from the list below

```

Model No	Model Type	Parameters
1	MobileNetV3-small	1.0 M
2	MobileNetV3-large	3.3 M
3	EfficientNetV2B0	6.3 M
4	EfficientNetV2B1	7.3 M
5	EfficientNetV2B2	9.1 M
6	InceptionResNetV2	54.7 M

Downloading data from https://storage.googleapis.com/tensorflow/keras-application/s/efficientnet_v2/efficientnetv2-b2_notop.h5

35839040/35839040 [=====] - 0s 0us/step

Created EfficientNetV2 B2 model with initial learning rate set to 0.001

enter the number of epochs to run initially. After these epochs complete you can enter H to halt training

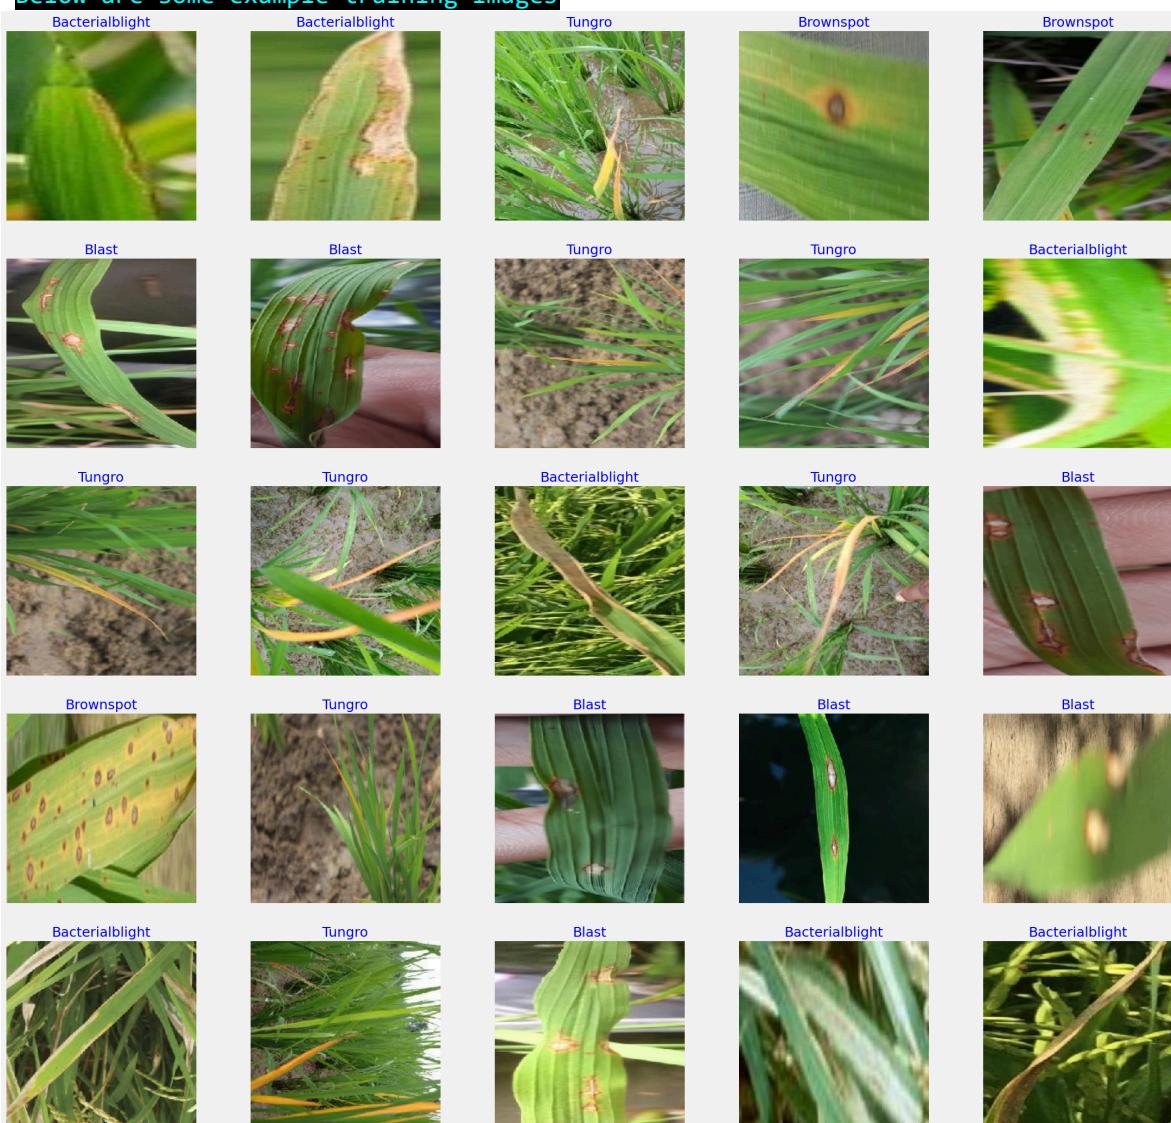
or enter an integer for how many more epochs to run then be asked again

Found 3840 validated image filenames belonging to 4 classes. for train generator

Found 480 validated image filenames belonging to 4 classes. for valid generator

Found 480 validated image filenames belonging to 4 classes. for test g enerator

Below are some example training images



Training will proceed until epoch 5 then you will be asked to

enter H to halt training or enter an integer for how many more epochs to run then be asked again

learning rate will be automatically adjusted during training

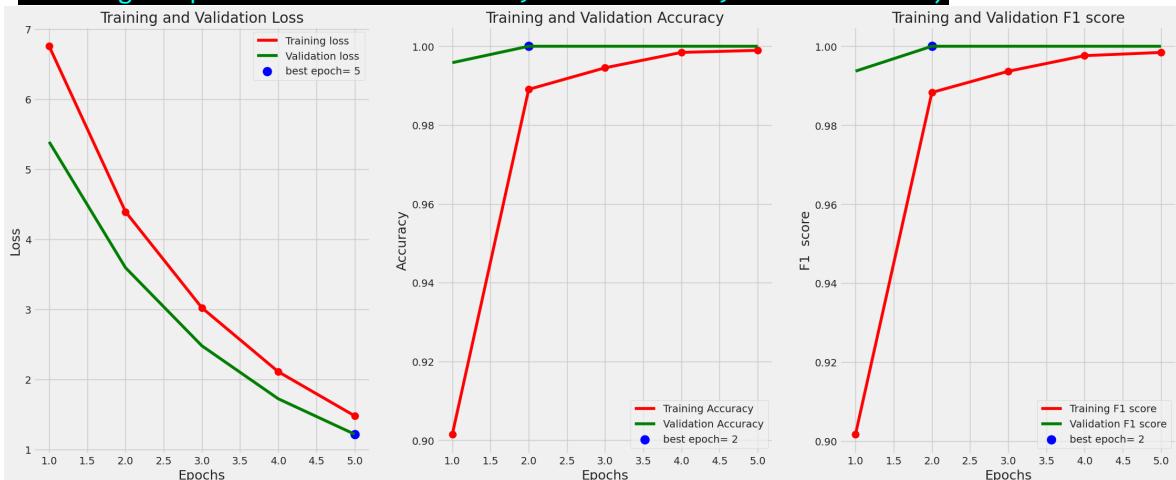
2024-01-01 12:24:08.921426: E tensorflow/core/grappler/optimizers/meta_optimizer.cc:954] layout failed: INVALID_ARGUMENT: Size of values 0 does not match size of permutation 4 @ fanin shape inmodel_3/block1b_drop/dropout/SelectV2-2-TransposeNHWCtoNCHW-LayoutOptimizer

Epoch	Train Duration in Seconds	Train Loss	Train Accuracy	Valid Loss	Valid Accuracy	V_Loss %	Learning Rate	Next LR	D
1	89.09	6.7565	90.16	5.3913	99.58	0.00	0.001000	0.001000	
2	32.99	4.3933	98.91	3.5943	100.00	33.33	0.001000	0.001000	
3	33.04	3.0224	99.45	2.4798	100.00	31.01	0.001000	0.001000	
4	33.02	2.1098	99.84	1.7236	100.00	30.49	0.001000	0.001000	
5	33.22	1.4799	99.90	1.2177	100.00	29.35	0.001000	0.001000	

Enter H to end training or an integer for the number of additional epochs to run then ask again
you entered h, Training halted on epoch 5 due to user input

loading model with weights from epoch 5

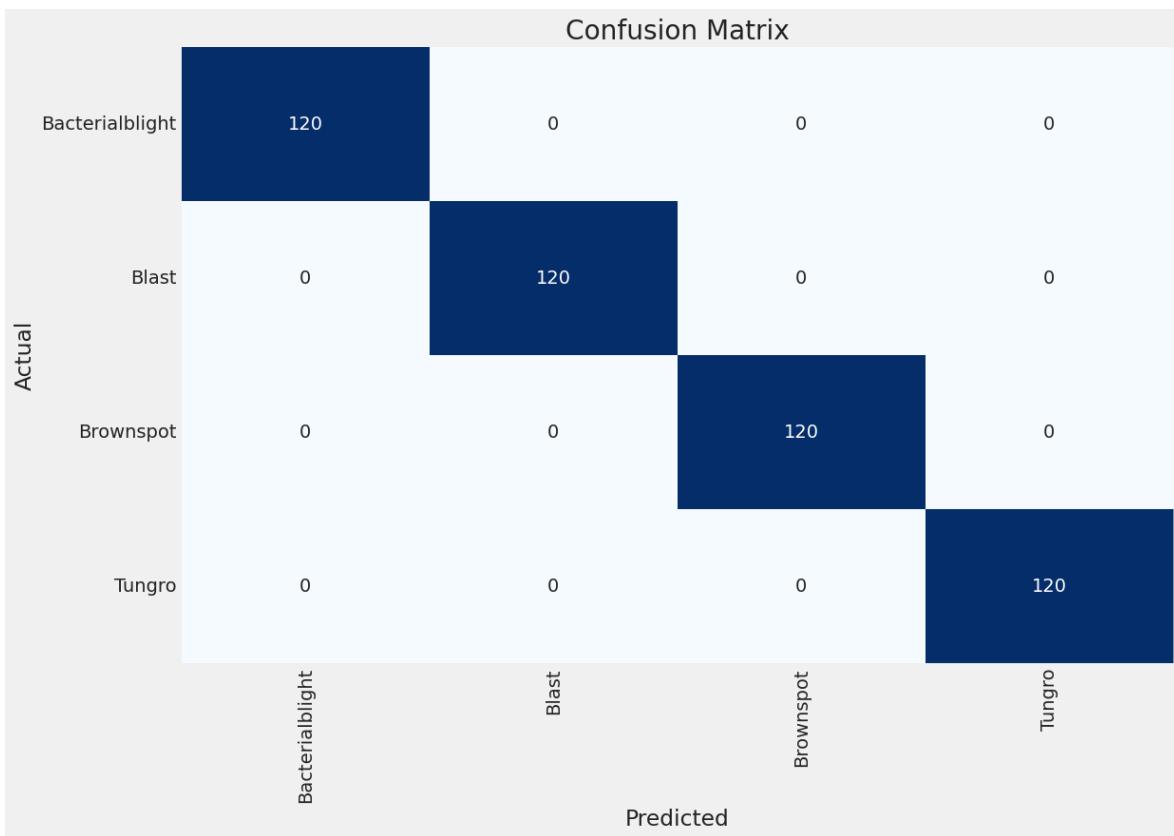
training elapsed time was 0.0 hours, 5.0 minutes, 31.09 seconds)



To save the training data to a csv file enter the name for the csv file or press enter to not save the data

6/6 [=====] - 4s 181ms/step

There were 0 errors in 480 tests for an accuracy of 100.00 and an F1 score of 100.00

**Classification Report:**

	precision	recall	f1-score	support
Bacterialblight	1.0000	1.0000	1.0000	120
Blast	1.0000	1.0000	1.0000	120
Brownspot	1.0000	1.0000	1.0000	120
Tungro	1.0000	1.0000	1.0000	120
accuracy			1.0000	480
macro avg	1.0000	1.0000	1.0000	480
weighted avg	1.0000	1.0000	1.0000	480

```
your trained model will be saved to directory /kaggle/working/ enter a subject for the saved model
model was saved as /kaggle/working/EfficientNetB2-4-(224 X 224)- 100.00.h5
model save nomenclature is directory/subject-number of classes- (img_height, img_width)- F1score.h5
enter R to rerun the classifier or press enter to quit
process completed
```

Model Evaluation

4 models were run with the same set of input image parameters. A limiter was set so that each class was limited to 1200 images. An image size of 224 X 224 was used. Since all classes have more than 1200 images the limited dataset is balanced and no image augmentation was required. Batch size was set to 40. The first model run was an EfficientNetB0 model. This model converged very quickly achieving a validation accuracy of 100% in

just 2 epochs. After 10 epochs the model achieved an F1 score of 100% with a total training time of 4 min and 19 seconds. The 2nd model run was a MobileNetV3 small model. This model converged slowly. After 10 epochs it achieved a validation accuracy of 91.7% After 15 epochs it achieved a validation accuracy of 100% The model was run for a total of 20 epochs and achieved an F1 score of 100% with a total run time of 4 minutes and 19 seconds. The 3rd model run was an MobileNetV3 large model. This model was run for 15 epochs and achieved an F1 score of 100% with a total run time of 4 minutes and 34 seconds. The final model run was an EfficientNetB2 model. This model converged very quickly. It was run for 5 epochs abd achieved an F1 score of 100% with a run time of 5 minutes and 31 seconds.